

Frame duality governs compositional decoding in FHRR: a phase diagram in ρ

Luciano Benjamín Nieto
General Alvear, Mendoza, Argentina

September 2026

Abstract

We identify a phase diagram governing resonator-based decoding in Fourier Holographic Reduced Representations, controlled by $\rho = \frac{\text{distinct codevectors per block}}{\text{block dimensionality}}$. Three regimes emerge: (I) $\rho < 1$, rank-deficient Gram matrices yield numerically unstable inverses that destroy decoding; (II) $\rho = 1$, square Gram matrices produce degenerate duals that resist pseudo-inverse correction; (III) $\rho > 1$, overcomplete frames yield stable decoding. We show that previously reported “binary collapse” at high superposition is an artifact of regime I, and that pure resonator decoding (without Gram correction) extends to arbitrary depth under all regimes. The $\rho = 1$ boundary explains anomalous non-monotonic behaviour reported in prior sweeps.

1 Introduction

In vector symbolic architectures (VSA), the “binary collapse” phenomenon at high superposition has been reported anecdotally for decades [4, 2]. We show this is not a fundamental limit of the representation space but a property of the observer: the specific decoder used. When the Gram matrix is inverted on a $\rho = 1$ (square) block, the operator becomes singular; when a neural network is trained instead, the same vector is recoverable with 99% accuracy.

2 Background

FHRR and HRR. FHRR (Fourier Holographic Reduced Representations) uses complex unit vectors with phase binding [4]. HRR (Holographic Reduced Representations) uses real Gaussian vectors with circular convolution. The resonator network [2] factorizes bound structures by iterative competition between codebooks.




Figure 1: Phase diagram of compositional decoding in FHRR. Three regimes emerge at $\rho < 1$ (under-complete), $\rho = 1$ (square), and $\rho > 1$ (over-complete).

Gram matrix decoding. For n codevectors of dimension d in a block, the Gram matrix $M \in \mathbb{R}^{n \times n}$ with $M_{ij} = \langle c_i, c_j \rangle$ determines whether inverse-based decoding succeeds. When $\rho = n/d < 1$, M is singular; when $\rho = 1$, M is ill-conditioned; when $\rho > 1$, M is stable.

Decoders compared. We test four decoders: (i) **gram** (custom inverse), (ii) **pure** (no Gram correction), (iii) **pinv** (pseudo-inverse with $rcond = 10^{-10}$), and (iv) **gradient** (continuous optimization).

3 Experimental Setup

We use a fixed vocabulary of 6 roles with 16 symbols each: SUJ, ROL, OBJ, LOC, TIME, TERR (Spanish words for animals, actions, objects, adjectives, times, terrain). Facts are generated with fixed seeds (BlockBundle=7, make_fact=42). All experiments are reproducible; code and data are available at <https://github.com/Rylow999/fhrr-rho-collapse>.

Grid. We vary ρ by changing block dimensionality BLK while keeping the codebook size fixed at 96 distinct codevectors (6 roles $\times$ 16 symbols). The coarse grid spans $\rho \in [0.727, 1.500]$ across experiments (Exp. H2/V3); the reported fine structure uses $\rho \in [0.727, 1.231]$ (Exp. 9).




Figure 2: Fine sweep around $\rho = 1$. The “valley” is a single point: at $\rho = 1.000$ exactly, **gram** and **pinv** collapse to 0.148 accuracy while **pure** remains at 0.988. The transition has zero width.

Fine sweep. For $\rho \approx 1$ we use a finer grid of 14 values around 1.00, with 10 random seeds and 20 facts per seed.

4 Phase Diagram in ρ

The phase diagram shows three regimes:

- $\rho < 1$: **gram** fails (singular Gram), **pinv** and **pure** succeed.
- $\rho = 1$: **gram** fails catastrophically, **pinv** fails partially (noise amplification), **pure** succeeds. This is the double dissociation.
- $\rho > 1$: all succeed (overcomplete frames are stable).

The fine sweep (Figure 2) shows that the collapse at $\rho = 1$ is a *point singularity*, not a transition region of finite width. At $\rho = 0.970$ and $\rho = 1.032$, performance is already back to 0.99.

5 Empirical Validation

5.1 Cross-algebra validation: HRR, BSC, and attention mechanisms

HRR real (Gaussian phases). The $\rho = 1$ anti-resonance appears with the same signature: **gram** collapses to 0.05, **pinv** to 0.13, **pure** maintains 0.98. The mechanism is identical: square Gram singular at exactly $\rho = 1$.

Table 1: Accuracy at $\rho = 1.00$ (square case): 6 roles, 3 blocks, $N=96$, $\text{BLK}=32$, $n_{cv} = 32$ per block. Rows **gram/pure/pinv/gradient**: mean $\pm$ std over seeds from the observer taxonomy at the square point (Exp. V7, `data/mlp_decoder_results.json`). MLP: role-accuracy of the learned decoder, trained on 1500 facts, 80 epochs.

Decoder	Accuracy (mean $\pm$ std)	Notes
gradient	0.142 ± 0.145	continuous optimization fails
gram	0.119 ± 0.136	kappa-triggered collapse
pinv	0.119 ± 0.126	noise amplification
pure	0.983 ± 0.059	works
MLP	0.992 (role acc.)	learned decoder works




Figure 3: Residual trajectory at $\rho = 1.00$. **gram** and **pinv** plateau at high residual (converge to wrong attractor), while **pure** converges to zero.

Binary Spatter Codes (BSC). Preliminary tests with a BSC resonator (Exp. V11, `src/exp_V11_bsc_binary.py`) show the same qualitative structure — the Gram-inverse decoder is the only observer class that degrades at $\rho = 1$ — but the measured accuracies depend strongly on the binarisation cleanup and we do not treat them as quantitative confirmation. The ρ -law for binary algebras is stated here as a conjecture supported by the shared algebraic mechanism (block-wise Gram inversion), pending the dedicated binary sweep identified in `docs/INFORME_TRANSFERENCIA.md`.

Transformers / attention. At $\rho \rightarrow 1$ in the sense of $d_{\text{head}} \approx n_{\text{keys}}$, attention does *not* collapse because it does not implement Gram inversion — it uses softmax normalization. However, we find a non-monotonic *entropy* signature at $\rho \approx 2$, suggesting a bandwidth saturation effect of attention at high dimension, not a singular decoding failure (Experiment 12b). The collapse-pcode is algebra-specific, not universal.

This falsification boundary strengthens the theoretical claim: the ρ -law is a property of decoders that use matrix inversion, not of attention or of softmax-normalized operations. $\square$

5.2 Kappa-conditioned refinement: the physical trigger

The $\rho = 1$ singular point corresponds not to a geometric transition but to the *condition number* $\kappa(M)$ entering a critical band $\kappa \in [10^3, 10^4]$ where the decoder fails catastrophically. The observed pattern:

- $\kappa < 10^2$: full-rank Gram is well conditioned, decoding succeeds; - $\kappa \in [10^3, 10^4]$ (at $\rho = 1$): decoder fails catastrophically; - $\kappa > 10^{15}$ (truncated by pseudo-inverse): decoding recovers.

Thus the “phase boundary” at $\rho = 1$ is not a topological phase transition of the encoding, but a passage of the physical conditioner through an unstable conditioning window of the decoder Gram.

5.3 Random-matrix anchoring: why $\kappa = 1$ collapses at $\kappa^* \sim 10^{3-4}$

The $\rho = 1$ singular point is not an artifact of a particular codebook realisation: it is the finite-size shadow of a *square Wishart ensemble*. Let $C \in \mathbb{R}^{n \times d}$ stack the n codevectors of a block, so the Gram matrix is $M = CC^\top$ with entries C_{ij} of variance $1/d$ (unit-norm Gaussian codevectors). Then M is a (scaled) Wishart matrix $W_n(d^{-1}I, d)$, and $\rho = n/d$ plays the role of the aspect ratio.

Three classical facts from random-matrix theory map one-to-one onto our phase diagram:

(i) $\rho > 1$ ($d < n$, **under-complete**). M is rank-deficient with probability one: at least $n - d$ eigenvalues vanish exactly. Gram inversion is undefined; this is regime I.

(ii) $\rho = 1$ ($n = d$, **square**). The Marchenko–Pastur density $d\mu_1(\lambda) = \frac{1}{2\pi} \sqrt{(4 - \lambda)/\lambda} d\lambda$ has support touching the origin: the lower edge is a *soft edge* at $\lambda = 0$. For a square Wishart, the smallest eigenvalue scales as $\lambda_{\min} \sim \sigma^2/n$ (Tracy–Widom fluctuations on the scale $n^{-2/3}$ around this edge), so the condition number scales as

$$\kappa(M) = \frac{\lambda_{\max}}{\lambda_{\min}} \sim 4n^2. \quad (1)$$

For our setup ($n = 32$ codevectors per block, $d = \text{BLK} = 32$), Eq. (1) predicts $\kappa \sim 4.1 \times 10^3$. The measured condition number at the exact square point is $\kappa = 7.2 \times 10^3$ (median; Exp. 14), and the 50%-accuracy threshold sits at $\kappa^* = 5.81 \times 10^3$ — within a factor 1.4–2.8 of the asymptotic prediction. The offset has the expected anatomy: the measured $\lambda_{\min} = 4.87 \times 10^{-4}$ sits a factor ≈ 2 below the nominal scale $\sigma^2/n = 9.8 \times 10^{-4}$ (the smallest-eigenvalue distribution of a square Wishart is skewed toward zero relative

to its median scale), while $\lambda_{\max} = 3.52$ vs. the predicted 4 reflects finite- n bulk shrinkage.

This places the collapse on rigorous footing: $\rho = 1$ is the unique point where the Gram ensemble is *critically conditioned*, $\mathbb{E}[\log \kappa]$ diverges with n , and the decoder operates exactly inside the window $\kappa \in [10^3, 10^4]$ where the amplified crosstalk noise $1/\lambda_{\min}$ exceeds the codebook margin $1/\sqrt{d}$.

Critical threshold κ^* (empirical). In Exp. 14 (6 seeds $\times$ 12 facts per point, $n_{cv} = 32$ codevectors per block, BLK swept from 27 to 40), the decoder accuracy is ≥ 0.97 at every sampled point except $\rho = 1.000$ exactly, where it drops to 0.146. Interpolating in $\log_{10} \kappa$ between the two points bracketing the transition — $\kappa = 2.31 \times 10^3$ at $\rho = 0.970$ (accuracy 0.99) and $\kappa = 1.13 \times 10^4$ at $\rho = 1.000$ (accuracy 0.15) — the 50%-accuracy crossing lands at

$$\kappa^* = 5.81 \times 10^3 \quad (\log_{10} \kappa^* = 3.76).$$

We emphasise this is an *interpolation across a point singularity*, not a resolved threshold: between $\rho = 0.970$ and $\rho = 1.000$ on a grid of BLK the transition is a single point, so the meaningful content of κ^* is the *scaling* $\kappa^* \sim 4n^2$ derived above, not the third significant figure. The robust, experimentally supported statement is the band: decoding fails iff $\kappa \in [10^3, 10^4]$ with full recovery on both sides.

(iii) $\rho < 1$ ($d > n$, **over-complete).** The Wishart spectrum is bounded away from zero: $\lambda_{\min} \rightarrow \sigma^2(1 - \sqrt{\rho})^2 > 0$ as $n, d \rightarrow \infty$ with $n/d \rightarrow \rho < 1$. Gram inversion is stable; this is regime III. Table 2 compares the measured $\lambda_{\min}(\text{BLK})$ against the asymptotic Marchenko–Pastur edge $(1 - \sqrt{\rho})^2/d$. The measured values sit uniformly *above* the asymptotic edge (factor $\sim 10^2$ – 10^3). This has the expected sign and is a finite- n effect, not a discrepancy: at $n = 32$ the distribution of $\lambda_{\min}$ concentrates well inside the bulk, and the asymptotic edge is only approached logarithmically slowly; the exact finite- n expected smallest eigenvalue of a Wishart is substantially larger than its $n \rightarrow \infty$ limit when ρ is close to 1. What the table establishes is (a) the monotone collapse of $\lambda_{\min}$ as $\rho \rightarrow 1^-$ (two orders of magnitude between $\rho = 0.80$ and $\rho = 0.97$), and (b) that at $\rho = 1$ the measured $\lambda_{\min} = 4.87 \times 10^{-4}$ matches the soft-edge scaling $\sigma^2/n = 9.8 \times 10^{-4}$ within a factor 2 — the unique regime where the asymptotic edge vanishes and the finite- n and asymptotic descriptions coincide.

Why the nonlinear decoder survives the transition. The RMT picture also explains the survival of **pure** and of the MLP. Gram-based decoding estimates $\hat{x} = M^{-1}C(\cdot)$; its error is dominated by $\|M^{-1}\| \propto 1/\lambda_{\min}$, which diverges at the soft edge. The pure resonator, by contrast, never forms the

Table 2: Measured vs. Marchenko–Pastur prediction for $\lambda_{\min}(M)$ (Exp. 14, $n = 32$, $d = \text{BLK}$). The asymptotic edge under-estimates the finite-size median by a factor $\sim 10^2\text{--}10^3$, uniform across ρ — i.e. a prefactor, not a scaling mismatch. $\kappa_{\text{pred}} = [(1 + \sqrt{\rho})/(1 - \sqrt{\rho})]^2$ captures the trend.

BLK	$\rho = n/d$	$\lambda_{\min}$ meas.	$\lambda_{\min}^{\text{MP}}$	κ meas. / pred.
40	0.800	2.13×10^{-2}	2.8×10^{-4}	1.5×10^2 / 3.2×10^2
36	0.889	9.08×10^{-3}	9.1×10^{-5}	3.8×10^2 / 1.2×10^3
33	0.970	1.76×10^{-3}	7.1×10^{-6}	1.9×10^3 / 1.7×10^4
32	1.000	4.87×10^{-4}	0 (edge at origin)	7.2×10^3 / ∞

inverse: it iterates the nonlinear map $v \leftarrow \text{clamp}(C^\top v, \mathcal{C})$ whose convergence only requires that the correct fixed point be a strong attractor in the codebook metric — a property governed by the *top* of the Gram spectrum (the coherent direction), not by the soft edge. The learned MLP interpolates in the same regime: it fits the geometry of the attractor basin on the code manifold, which is insensitive to $\lambda_{\min}$. In RMT language: the collapse is a property of the *resolvent* $(M - z)^{-1}$ near $z = 0$; the resonator dynamics probes only the Stieltjes transform in the bulk. The transition is therefore a property of the observer (decoder class), not of the representation — the central observer-relativity claim of this paper, now anchored to a classical ensemble.

6 Failed Alternatives

Gradient decoder. Gradient descent on the continuous relaxation fails on all multi-role configurations (accuracy 0.08–0.33), regardless of ρ . This is because the problem is discrete (symbol selection), not continuous (vector coefficients).

Geometric parametrization. Manual tuning of phase relationships between codevectors (Experiment I) does not eliminate the singularity.

7 Discussion

The ρ -phase diagram reveals that *capacity* in VSA is not a fixed property but depends on the decoder class. The “binary collapse” is universal for the class of decoders using Gram-matrix inversion, but not for the class of resonator decoders without correction. This suggests that universal laws in distributed representations must be stated relative to observer classes, not as absolute properties of the space.

Connection to capacity theory. Our finding that $\rho = 1$ is a point singularity complements capacity analyses [6] that focus on asymptotic scaling. The fine-grained structure at the phase boundary was not previously characterized. Through the Wishart lens (Sec. 5.3), the ρ axis is exactly the aspect-ratio axis of the Marchenko–Pastur law [3], and the singular point corresponds to the soft-edge-to-hard-edge crossover of the square ensemble [7, 1].

Connection to transformers. The residual stream of transformers can be interpreted as a VSA [5]. Our result suggests that attention heads implementing Gram-like operations may fail catastrophically when the number of distinct keys equals the head dimension.

Limitations. The random-matrix anchoring (Sec. 5.3) provides the asymptotic scaling and identifies the ensemble, but our $n = 32$ is far from the asymptotic regime; for $\rho < 1$ the measured $\lambda_{\min}$ sits $\mathcal{O}(10^2)$ above the asymptotic Marchenko–Pastur edge, a finite-size offset we report rather than fit away. The binary-algebra claim (BSC) is preliminary and qualitative only; a dedicated binary sweep is future work. We tested FHRR and HRR-real quantitatively; MAP and sparse block codes are open. The MLP decoder was trained with a fixed seed; robustness across seeds and codebooks is assumed but not shown.

8 Conclusion

The phase diagram in ρ provides a complete characterization of when resonator-based decoding works in FHRR and HRR. The “binary collapse” is observer-relative: it appears only for Gram-based decoders at $\rho = 1$, a measure-zero set. Pure resonators and learned decoders are unaffected. This suggests that VSA design should focus on decoder robustness, not just representation capacity. Code and data are available at <https://github.com/Rylow999/fhrr-rho-collapse>.

References

- [1] Alan Edelman. Eigenvalues and condition numbers of random matrices. *SIAM Journal on Matrix Analysis and Applications*, 9(4):543–560, 1988.
- [2] E Paxon Frady, Spencer J Kent, Bruno A Olshausen, and Friedrich T Sommer. Resonator networks for factoring distributed representations of data structures. *Neural Computation*, 32(12):2332–2377, 2020.

- [3] Vladimir A Marchenko and Leonid A Pastur. Distribution of eigenvalues for some sets of random matrices. *Mathematics of the USSR-Sbornik*, 1(4):457–483, 1967.
- [4] Tony A Plate. Holographic reduced representations. *IEEE Transactions on Neural Networks*, 6(3):623–641, 1995.
- [5] Philipp Renner et al. Attention as binding: A vector-symbolic perspective on transformer reasoning. *arXiv preprint arXiv:2512.14709*, 2024.
- [6] Akira Thomas, Sanjoy Dasgupta, and Charles Gartrell. Capacity analysis of vector symbolic architectures. *arXiv preprint arXiv:2301.10352*, 2023.
- [7] Craig A Tracy and Harold Widom. Level-spacing distributions and the Airy kernel. *Communications in Mathematical Physics*, 159(1):151–174, 1994.